

Logistic regression with Amazon food review dataset

In [1]:

```
import pandas as pd
import numpy as np
import pickle
#cleaned_data = pickle.load(open('cleaned.p', 'rb'))

cleaned_data = pickle.load(open('clen.p', 'rb')) ##
print(cleaned_data.shape)
print(type(cleaned_data))

cleaned_data=cleaned_data.iloc[:100000]
cleaned_data.shape
```

```
(250000, 11)
<class 'pandas.core.frame.DataFrame'>
```

Out[1]:

```
(100000, 11)
```

In [2]:

```
#Sort the data with accordance to time
cleaned_data.sort_values('Time', inplace=True)
cleaned_data.head(3)
```

Out[2]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	
138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	0	positive	93
138683	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	2	positive	94
417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	0	positive	94

In [3]:

```
# Find if the Y is imbalanced.
#positive_negative = cleaned_data['Score']
#print (positive_negative.value_counts())
#positive_negative.shape# -----found the Y is imbalanced

def partition(x):
    if x == 'negative':
        return '0'
    return '1'
actualScore = cleaned_data['Score']
positiveNegative = actualScore.map(partition)
cleaned_data['Score'] = positiveNegative
```

In [4]:

```
#positive_negative= cleaned_data['Score'].map(lambda x: 1 if int (x) is 'positive' else 0)
#print (positive_negative.value_counts())
positiveNegative = cleaned_data['Score']
print (positiveNegative.value_counts())
positiveNegative.shape

cleaned_data.head(5)
```

```
1    87730
0    12270
Name: Score, dtype: int64
```

Out[4]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	
138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	0	1	93
138683	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	2	1	94
417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	0	1	94
346055	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1	2	1	94
417838	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0	0	1	94

In [5]:

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer

X_tra, X_tes, y_train, y_test = train_test_split(cleaned_data['CleanedText'].values, positiveNegative,
test_size=0.3, shuffle=False)

#Implementing BAG of words
tf_idf_vect = TfidfVectorizer(ngram_range=(1,0), dtype=float)
X_tf = tf_idf_vect.fit_transform(X_tra)

# Standerdising the data
norm = StandardScaler(with_mean = False)
X_train = norm.fit_transform(X_tf)

# tfidf test
X_tfte = tf_idf_vect.transform(X_tes)

# Standerdising the data
X_test = norm.transform(X_tfte)
''' Use transform for follwing function (vectorizer and standerdisation) after main function
and not applicable for actual models'''
```

```

and not applicable for actual models
#uni_gram = CountVectorizer(ngram_range=(1,1))
#X_BOW = uni_gram.fit_transform(X_tra)
# Standerdising the data
#X_train2 = StandardScaler(with_mean = False).fit_transform(X_BOW)

#vectorizer for test data
#X_BOW1 = uni_gram.transform(X_tes)
# Standerdising the data
#X_test2 = StandardScaler(with_mean = False).fit_transform(X_BOW1)

```

```

from sklearn.model_selection import TimeSeriesSplit
tscv = TimeSeriesSplit(n_splits=10)
for train, cv in tscv.split(X_train):
    print(X_train[train].shape, X_train[cv].shape)

```

```

(6370, 41598) (6363, 41598)
(12733, 41598) (6363, 41598)
(19096, 41598) (6363, 41598)
(25459, 41598) (6363, 41598)
(31822, 41598) (6363, 41598)
(38185, 41598) (6363, 41598)
(44548, 41598) (6363, 41598)
(50911, 41598) (6363, 41598)
(57274, 41598) (6363, 41598)
(63637, 41598) (6363, 41598)

```

In [9]:

```

#GridSearch with default l2 norm
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import TimeSeriesSplit
from sklearn.metrics import make_scorer
from sklearn.metrics import f1_score

f1 = make_scorer(f1_score, pos_label='1')
tscv = TimeSeriesSplit(n_splits=3)
param_grid = {'C':[10**-4, 10**-2, 10**0, 10**2, 10**4]}
log = LogisticRegression(class_weight='balanced')
gsv = GridSearchCV(log,param_grid,cv=tscv,scoring=f1)
gsv.fit(X_train,y_train)

print("Best HyperParameter: ",gsv.best_params_)
print("Best f1: %.2f%%"%(gsv.best_score_*100))

```

```

Best HyperParameter:  {'C': 0.0001}
Best f1: 94.16%

```

In [10]:

```

print(gsv.best_estimator_)

```

```

LogisticRegression(C=0.0001, class_weight='balanced', dual=False,
    fit_intercept=True, intercept_scaling=1, max_iter=100,
    multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
    solver='liblinear', tol=0.0001, verbose=0, warm_start=False)

```

In [15]:

```

print(gsv.best_estimator_)
print("Best HyperParameter: ",gsv.best_params_)
print("Best f1: %.2f%%"%(gsv.best_score_*100))
import matplotlib.pyplot as plt
x=[]
y=[]
for a in gsv.grid_scores_:
    x.append(a[0]['C'])
    y.append(a[1])
plt.xlim(-1,50)
plt.ylim(0.8,1)
plt.xlabel(r"$\ C = 1/\text{Lambda} \ $",fontsize=15)

```

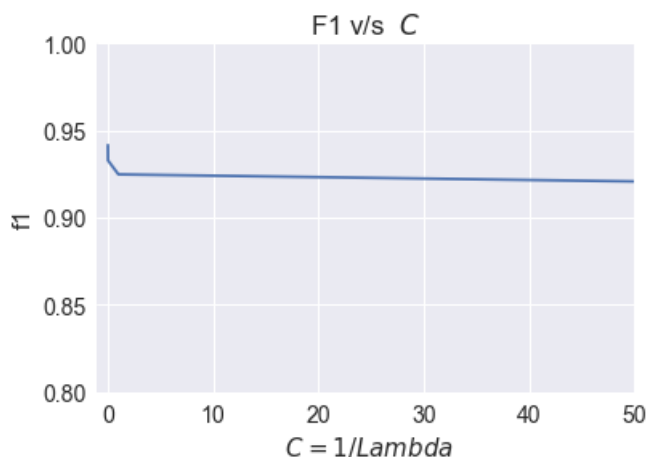
```
plt.ylabel("f1")
plt.title(r'F1 v/s  $C$ ')
plt.plot(x,y)
plt.show()
```

```
LogisticRegression(C=0.0001, class_weight='balanced', dual=False,
                    fit_intercept=True, intercept_scaling=1, max_iter=100,
                    multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
                    solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
```

Best HyperParameter: {'C': 0.0001}

Best f1: 94.16%

C:\Users\admin\Anaconda3\lib\site-packages\sklearn\model_selection_search.py:761:
 DeprecationWarning: The grid_scores_ attribute was deprecated in version 0.18 in favor of the more
 elaborate cv_results_ attribute. The grid_scores_ attribute will not be available from 0.20
 DeprecationWarning)



In [14]:

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.linear_model import LogisticRegression
import seaborn as sns

clf = LogisticRegression(C= 0.0001, penalty= 'l2')
clf.fit(X_train,y_train)
y_pred = clf.predict(X_train)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_train, y_pred)*100))
print("Non Zero weights:",np.count_nonzero(clf.coef_))
print("F1-Score on test set: %0.3f%%"%(f1_score(y_train, y_pred, pos_label='1')*100))
print (classification_report(y_train,y_pred))
result = confusion_matrix(y_train,y_pred)
print("Confusion Matrix of test set:\n [ [TN  FP]\n [FN TP] ]\n")
print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')
```

Accuracy on test set: 95.491%

Non Zero weights: 41598

F1-Score on test set: 97.507%

	precision	recall	f1-score	support
0	0.98	0.63	0.77	8172
1	0.95	1.00	0.98	61828
avg / total	0.96	0.95	0.95	70000


Confusion Matrix of test set:

```
[ [TN  FP]
  [FN TP] ]
```

```
[[ 5137  3035]
 [  121 61707]]
```

Out[14]:

<matplotlib.axes._subplots.AxesSubplot at 0xc7619e8>



In [16]:

```
clf = LogisticRegression(C=1, penalty='l2')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))
```

Non Zero weights: 41598

In [17]:

```
clf = LogisticRegression(C=1, penalty='l1')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))
```

Non Zero weights: 15981

In [18]:

```
clf = LogisticRegression(C=100, penalty='l2')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))
```

Non Zero weights: 41598

In [19]:

```
clf = LogisticRegression(C=100, penalty='l1')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))
```

Non Zero weights: 18136

In [20]:

```
#test
clf = LogisticRegression(C=1, penalty='l1');
model=clf.fit(X_train, y_train);
pred = (model.predict(X_test))
```

In [21]:

```
print(classification_report(y_test,pred))
result = confusion_matrix(y_test,pred)
print("Confusion Matrix of test set:\n [ [TN  FP]\n [FN TP] ]\n")
print(result)
sns.set(font_scale=1.4)#for label size
```

```

sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')

print("f1 on test set: %0.3f%%"%(f1_score(y_test, pred,pos_label='1')*100))
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, pred)*100))

```

	precision	recall	f1-score	support
0	0.55	0.55	0.55	4098
1	0.93	0.93	0.93	25902
avg / total	0.88	0.88	0.88	30000

Confusion Matrix of test set:

```


[ [TN FP]
  [FN TP] ]

[[ 2254  1844]
 [ 1864 24038]]

```

f1 on test set: 92.839%

Accuracy on test set: 87.640%



In [42]:

```

# Feature importance

all_feat = tf_idf_vect.get_feature_names()

feature_coefs = pd.DataFrame(
    data = list(zip(all_feat, model.coef_[0])),
    columns = ['-ve feature', 'coef'])

feature_imp = feature_coefs.sort_values(by='coef', ascending=True)
feature_imp.head(8)

```

Out[42]:

	-ve feature	coef
40796	worst	-0.494986
10305	disappoint	-0.460843
2196	author	-0.416992
35729	tast	-0.387144
2287	aw	-0.364683
14110	fragment	-0.353957
16501	hard	-0.322843
34073	stale	-0.317204

In [43]:

```

feature_coefs1 = pd.DataFrame(
    columns = ['+ve feature', 'coef'],
    data = list(zip(all_feat, model.coef_[0])))

```

```
feature_imp = feature_coefs1.sort_values(by='coef', ascending=False)
feature_imp.head(8)
```

Out[43]:

	+ve feature	coef
15809	great	1.471587
3318	best	1.246960
21214	love	1.134374
41378	yum	0.993332
20488	leg	0.947100
16758	heartili	0.886589
41391	yummi	0.837720
9628	delici	0.832593

In [178]:

```
# Perturbation tech
from scipy.sparse import find

clf = LogisticRegression(C=0.0001, penalty='l2')
clf.fit(X_train, y_train)
weights_before = clf.coef_

epsilon = np.random.normal(scale=0.1)
print ('noise adding to weight vector is = {0}'.format(epsilon))
X_train1 = X_train
a,b = np.nonzero(X_train1)
X_train1[a,b] = X_train[a,b]+epsilon
clf.fit(X_train1,y_train)
weights_after = clf.coef_
print ('weights_after adding noise = {0}'.format(weights_after))
```

```
noise adding to weight vector is = -0.05678220906943951
weights_after adding noise = [[ 1.40735563e-03  1.17262635e-03  1.45049363e-03 ... -5.78981902e-03
 2.22507124e-04  5.26017521e-05]]
```

In [179]:

```
weights_diff = abs(weights_before - weights_after)
per_weight_diff1=[]
per_weight_diff1 = weights_diff-0.3*abs (weights_before)
(per_weight_diff1[per_weight_diff1>0].size)/weights_before.size*100
```

Out[179]:

0.2497055359246172

In [181]:

```
import numpy
weights_diff.shape
value1 = numpy.where( weights_diff > 0.3 )# numpy for above 30%
print(value1)
weights_diff[ numpy.where( weights_diff > 0.01 ) ] # find the exact value above 0.01
```

```
(array([], dtype=int64), array([], dtype=int64))
```

Out[181]:

array([], dtype=float64)

Random search model

In [22]:

```
# RandomsearchCV with default l2 norm
#from sklearn.model_selection import TimeSeriesSplit
#print(X_train[train].shape, X_train[cv].shape)
from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import RandomizedSearchCV
from sklearn.linear_model import LogisticRegression

f1 = make_scorer(f1_score, pos_label='1')
tscv = TimeSeriesSplit(n_splits=3)
#clf = LogisticRegression(class_weight='balanced',penalty='l2')
#params we need to try on classifier
param_grid = {'C':[10**-4, 10**-2, 10**0, 10**2, 10**4]}

rsv = RandomizedSearchCV(LogisticRegression(class_weight='balanced',penalty='l2'),
param_grid,n_iter=5,cv=tscv,scoring=f1)
rsv.fit(X_train,y_train)

#savetofile(gsv,"Log Reg/gsv_uni")
print("Best HyperParameter: ",rsv.best_params_)
print("Best F1 score: %.2f%%"%(rsv.best_score_*100))
```

Best HyperParameter: {'C': 0.0001}
Best F1 score: 94.16%

In [23]:

```
clf = LogisticRegression(C= 0.0001, penalty= 'l2')
clf.fit(X_train,y_train)
y_pred = clf.predict(X_train)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_train, y_pred)*100))
print("Non Zero weights:",np.count_nonzero(clf.coef_))
print("F1-Score on test set: %0.3f%%"%(f1_score(y_train, y_pred, pos_label='1')*100))
print (classification_report(y_train,y_pred))
result = confusion_matrix(y_train,y_pred)
print("Confusion Matrix of test set:\n [ [TN  FP]\n [FN TP] ]\n")
print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')
```

Accuracy on test set: 95.491%
Non Zero weights: 41598
F1-Score on test set: 97.507%


	precision	recall	f1-score	support
0	0.98	0.63	0.77	8172
1	0.95	1.00	0.98	61828
avg / total	0.96	0.95	0.95	70000

Confusion Matrix of test set:
[[TN FP]
[FN TP]]

```
[[ 5137  3035]
 [ 121 61707]]
```

Out[23]:

<matplotlib.axes._subplots.AxesSubplot at 0xd8eaac8>





In [24]:

```
clf = LogisticRegression(C=1, penalty='l2')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))
```

Non Zero weights: 41598

In [25]:

```
clf = LogisticRegression(C=1, penalty='l1')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))
```

Non Zero weights: 15942

In [26]:

```
clf = LogisticRegression(C=100, penalty='l2')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))
```

Non Zero weights: 41598

In [27]:

```
clf = LogisticRegression(C=100, penalty='l1')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))
```

Non Zero weights: 18125

In [28]:

```
clf = LogisticRegression(C=0.001, penalty='l2');
model=clf.fit(X_train, y_train);
pred = (model.predict(X_test))
```

In [29]:

```
print(classification_report(y_test,pred))
result = confusion_matrix(y_test,pred)
print("Confusion Matrix of test set:\n [ [TN FP]\n [FN TP] ]\n")
print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')

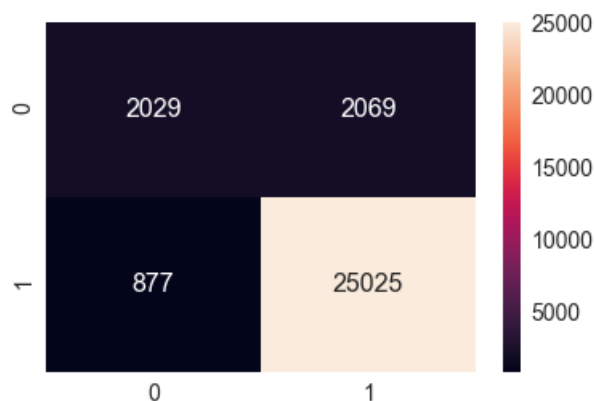
print("f1 on test set: %0.3f%%"%(f1_score(y_test, pred, pos_label='1')*100))
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, pred)*100))
```

	precision	recall	f1-score	support
0	0.70	0.50	0.58	4098
1	0.92	0.97	0.94	25902
avg / total	0.89	0.90	0.89	30000

Confusion Matrix of test set:

	TN	FP
FN	TP	

```
[[ 2029 2069]
 [ 877 25025]]
f1 on test set: 94.441%
Accuracy on test set: 90.180%
```



LR with Grid search CV	LR with Random search CV
1. Hyper parameter is 0.001	Hyper parameter is 0.001
2. F1 score of Train data = 97.5%	F1 score of Train data = 97.5%
3. F1 score of Test data	F1 score of Test data
92.85%	94.4%

BOW Logistic regression

Grid search

In [30]:

```
#count_vect = CountVectorizer(1,1)
#count_vect.fit_transform(final['CleanedText'].values)

#Implementing BAG of words
bi_gram = CountVectorizer(1,0)
X_BOW = bi_gram.fit_transform(X_tra)
# Standerdising the data
norm = StandardScaler(with_mean = False)
X_train = norm.fit_transform(X_BOW)
#X_train = preprocessing.normalize(X_BOW)

#vectorizer for test data
X_BOW1 = bi_gram.transform(X_tes)
# Standerdising the data
X_test = norm.transform(X_BOW1)

from sklearn.model_selection import TimeSeriesSplit

for train, cv in tscv.split(X_train):
    print(X_train[train].shape, X_train[cv].shape)
```

```
C:\Users\admin\Anaconda3\lib\site-packages\sklearn\utils\validation.py:475: DataConversionWarning:
Data with input dtype int64 was converted to float64 by StandardScaler.
    warnings.warn(msg, DataConversionWarning)
C:\Users\admin\Anaconda3\lib\site-packages\sklearn\utils\validation.py:475: DataConversionWarning:
Data with input dtype int64 was converted to float64 by StandardScaler.
    warnings.warn(msg, DataConversionWarning)
```

```
(17500, 41598) (17500, 41598)
(35000, 41598) (17500, 41598)
(52500, 41598) (17500, 41598)
```

```
C:\Users\admin\Anaconda3\lib\site-packages\sklearn\utils\validation.py:475: DataConversionWarning:
Data with input dtype int64 was converted to float64 by StandardScaler.
    warnings.warn(msg, DataConversionWarning)
```

In [31]:

```
f1 = make_scorer(f1_score, pos_label='1')
tscv = TimeSeriesSplit(n_splits=3)
param_grid = {'C':[10**-4, 10**-2, 10**0, 10**2, 10**4]}
log = LogisticRegression(class_weight='balanced')
gsv = GridSearchCV(log,param_grid,cv=tscv,scoring=f1)
gsv.fit(X_train,y_train)

print("Best HyperParameter: ",gsv.best_params_)
print("Best f1: %.2f%%"%(gsv.best_score_*100))
```

Best HyperParameter: {'C': 0.0001}
Best f1: 94.61%

In [32]:

```
print(gsv.best_estimator_)
```


```
LogisticRegression(C=0.0001, class_weight='balanced', dual=False,
                    fit_intercept=True, intercept_scaling=1, max_iter=100,
                    multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
                    solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
```

In [33]:

```
print("Best HyperParameter: ",gsv.best_params_)
print("Best f1: %.2f%%"%(gsv.best_score_*100))
import matplotlib.pyplot as plt
x=[]
y=[]
for a in gsv.grid_scores_:
    x.append(a[0]['C'])
    y.append(a[1])
plt.xlim(-1,50)
plt.ylim(0.8,1)
plt.xlabel(r"$\ C = 1/\text{Lambda} \$",fontsize=15)
plt.ylabel("f1")
plt.title(r'$F1$ v/s $C$')
plt.plot(x,y)
plt.show()
```

Best HyperParameter: {'C': 0.0001}
Best f1: 94.61%

C:\Users\admin\Anaconda3\lib\site-packages\sklearn\model_selection_search.py:761:
DeprecationWarning: The grid_scores_ attribute was deprecated in version 0.18 in favor of the more
elaborate cv_results_ attribute. The grid_scores_ attribute will not be available from 0.20
DeprecationWarning)



In [34]:

```

clf = LogisticRegression(C= 0.0001, penalty= 'l2')
clf.fit(X_train,y_train)
y_pred = clf.predict(X_train)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_train, y_pred)*100))
print("Non Zero weights:",np.count_nonzero(clf.coef_))
print("F1-Score on test set: %0.3f%%"%(f1_score(y_train, y_pred, pos_label='1')*100))
print(classification_report(y_train,y_pred))
result = confusion_matrix(y_train,y_pred)
print("Confusion Matrix of test set:\n [ [TN  FP]\n [FN TP] ]\n")
print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')

```

Accuracy on test set: 95.297%

Non Zero weights: 41598

F1-Score on test set: 97.403%

	precision	recall	f1-score	support
0	0.98	0.61	0.75	8172
1	0.95	1.00	0.97	61828
avg / total	0.95	0.95	0.95	70000

Confusion Matrix of test set:

```


[ [TN  FP]
  [FN TP] ]

[[ 4984  3188]
 [   104 61724]]

```

Out[34]:

<matplotlib.axes._subplots.AxesSubplot at 0xdc78e10>



In [38]:

```

clf = LogisticRegression(C=1, penalty='l2')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))

```

Non Zero weights: 41598

In [39]:

```

clf = LogisticRegression(C=1, penalty='l1')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))

```

Non Zero weights: 13909

In [40]:

```

clf = LogisticRegression(C=100, penalty='l2')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))

```

Non Zero weights: 41598

Non Zero weights: 41598

In [21]:

```
clf = LogisticRegression(C=100, penalty='l1')
model=clf.fit(X_train, y_train)
print("Non Zero weights:", np.count_nonzero(clf.coef_))
```

Non Zero weights: 18201

In [22]:

```
#test
clf = LogisticRegression(C=1, penalty='l1');
model=clf.fit(X_train, y_train);
pred = (model.predict(X_test))
```

In [23]:

```
print(classification_report(y_test, pred))
result = confusion_matrix(y_test, pred)
print("Confusion Matrix of test set:\n [ [TN FP]\n [FN TP] ]\n")
print(result)
sns.set(font_scale=1.4) #for label size
sns.heatmap(result, annot=True, annot_kws={"size": 16}, fmt='g')

print("f1 on test set: %0.3f%%"%(f1_score(y_test, pred, pos_label='1')*100))
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, pred)*100))
```

	precision	recall	f1-score	support
0	0.55	0.55	0.55	4098
1	0.93	0.93	0.93	25902
avg / total	0.88	0.88	0.88	30000


Confusion Matrix of test set:

```
[ [TN FP]
  [FN TP] ]
```

```
[[ 2261 1837]
 [ 1858 24044]]
```

f1 on test set: 92.864%

Accuracy on test set: 87.683%



Random search model

In [12]:

```
from sklearn.metrics import make_scorer
from sklearn.metrics import f1_score
from sklearn.model_selection import RandomizedSearchCV
from sklearn.linear_model import LogisticRegression

f1 = make_scorer(f1_score, pos_label='1')
```

```

tscv = TimeSeriesSplit(n_splits=3)
#clf = LogisticRegression(class_weight='balanced',penalty='l2')
#params we need to try on classifier
param_grid = {'C':[10**-4, 10**-2, 10**0, 10**2, 10**4]}

rsv = RandomizedSearchCV(LogisticRegression(class_weight='balanced',penalty='l2'),
param_grid,n_iter=5,cv=tscv,scoring=f1)
rsv.fit(X_train,y_train)

#savetofile(gsv,"Log Reg/gsv_uni")
print("Best HyperParameter: ",rsv.best_params_)
print("Best F1 score: %.2f%%"%(rsv.best_score_*100))

```

Best HyperParameter: {'C': 0.0001}
Best F1 score: 94.16%

In [14]:

```

from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.linear_model import LogisticRegression
import seaborn as sns

clf = LogisticRegression(C= 0.0001, penalty= 'l2')
clf.fit(X_train,y_train)
y_pred = clf.predict(X_train)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_train, y_pred)*100))
print("Non Zero weights:",np.count_nonzero(clf.coef_))
print("F1-Score on test set: %0.3f%%"%(f1_score(y_train, y_pred, pos_label='1')*100))
print (classification_report(y_train,y_pred))
result = confusion_matrix(y_train,y_pred)
print("Confusion Matrix of test set:\n [ [TN  FP]\n [FN TP] ]\n")
print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')

```

Accuracy on test set: 95.491%
Non Zero weights: 41598
F1-Score on test set: 97.507%

	precision	recall	f1-score	support
0	0.98	0.63	0.77	8172
1	0.95	1.00	0.98	61828
avg / total	0.96	0.95	0.95	70000

Confusion Matrix of test set:

```

[ [TN  FP]
  [FN TP] ]

[[ 5137  3035]
 [  121 61707]]

```

Out[14]:

<matplotlib.axes._subplots.AxesSubplot at 0xc8b03c8>

In [15]:

```

clf = LogisticRegression(C=1, penalty='l2')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))

```

Non Zero weights: 41598

In [16]:

```

clf = LogisticRegression(C=100, penalty='l1')

```

```
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))
```

Non Zero weights: 18155

In [17]:

```
clf = LogisticRegression(C=100, penalty='l2')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))
```

Non Zero weights: 41598

In [18]:

```
clf = LogisticRegression(C=1, penalty='l2')
model=clf.fit(X_train, y_train)
print("Non Zero weights:",np.count_nonzero(clf.coef_))
```

Non Zero weights: 41598

In [19]:

```
#test
clf = LogisticRegression(C=1, penalty='l1');
model=clf.fit(X_train, y_train);
pred = (model.predict(X_test))
```

In [20]:

```
print(classification_report(y_test,pred))
result = confusion_matrix(y_test,pred)
print("Confusion Matrix of test set:\n [ [TN FP]\n [FN TP] ]\n")
print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')

print("f1 on test set: %0.3f%%"%(f1_score(y_test, pred, pos_label='1')*100))
print("Accuracy on test set: %0.3f%%"%(accuracy_score(y_test, pred)*100))
```

	precision	recall	f1-score	support
0	0.55	0.55	0.55	4098
1	0.93	0.93	0.93	25902
avg / total	0.88	0.88	0.88	30000

Confusion Matrix of test set:

```
[ [TN FP]
 [FN TP] ]
```

```
[[ 2252 1846]
 [ 1865 24037]]
```

f1 on test set: 92.834%

Accuracy on test set: 87.630%



LR with Grid search CV	LR with Random search CV
1. Hyper parameter is 0.001	Hyper parameter is 0.001
2. F1 score of Train data = 97.5%	F1 score of Train data = 97.5%
3. F1 score of Test data	F1 score of Test data
92.85%	92.83%

Logistic Regression with average word2vec

In [74]:

```
import pandas as pd
import numpy as np
import pickle

cleaned_data = pickle.load(open('drive/My Drive/Colab Notebooks/new.p', 'rb')) ##
print(cleaned_data.shape)
print(type(cleaned_data))

cleaned_data=cleaned_data.iloc[:100000]
cleaned_data.shape
```

```
(250000, 11)
<class 'pandas.core.frame.DataFrame'>
```

Out[74]:

```
(100000, 11)
```

In [75]:

```
#Sort the data with accordance to time
cleaned_data.sort_values('Time', inplace=True)
cleaned_data.head(3)
```

Out[75]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Star
138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	0	positive	939340800	edi
138683	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	2	positive	940809600	grei sp
417839	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	0	positive	944092800	Ente

In [76]:

```
!pip install gensim
```

```
Requirement already satisfied: gensim in /usr/local/lib/python3.6/dist-packages (3.6.0)
Requirement already satisfied: scipy>=0.18.1 in /usr/local/lib/python3.6/dist-packages (from
gensim) (0.19.1)
Requirement already satisfied: six>=1.5.0 in /usr/local/lib/python3.6/dist-packages (from gensim)
(1.11.0)
Requirement already satisfied: numpy>=1.11.3 in /usr/local/lib/python3.6/dist-packages (from
gensim) (1.14.6)
Requirement already satisfied: smart-open>=1.2.1 in /usr/local/lib/python3.6/dist-packages (from
gensim) (1.7.1)
Requirement already satisfied: boto3 in /usr/local/lib/python3.6/dist-packages (from smart-
open>=1.2.1->gensim) (1.9.23)
Requirement already satisfied: requests in /usr/local/lib/python3.6/dist-packages (from smart-
open>=1.2.1->gensim) (2.18.4)
Requirement already satisfied: boto>=2.32 in /usr/local/lib/python3.6/dist-packages (from smart-
open>=1.2.1->gensim) (2.49.0)
```

```
Requirement already satisfied: bz2file in /usr/local/lib/python3.6/dist-packages (from smart-open>=1.2.1->gensim) (0.98)
Requirement already satisfied: boto3<1.13.0,>=1.12.23 in /usr/local/lib/python3.6/dist-packages (from boto3->smart-open>=1.2.1->gensim) (1.12.23)
Requirement already satisfied: jmespath<1.0.0,>=0.7.1 in /usr/local/lib/python3.6/dist-packages (from boto3->smart-open>=1.2.1->gensim) (0.9.3)
Requirement already satisfied: s3transfer<0.2.0,>=0.1.10 in /usr/local/lib/python3.6/dist-packages (from boto3->smart-open>=1.2.1->gensim) (0.1.13)
Requirement already satisfied: chardet<3.1.0,>=3.0.2 in /usr/local/lib/python3.6/dist-packages (from requests->smart-open>=1.2.1->gensim) (3.0.4)
Requirement already satisfied: urllib3<1.23,>=1.21.1 in /usr/local/lib/python3.6/dist-packages (from requests->smart-open>=1.2.1->gensim) (1.22)
Requirement already satisfied: idna<2.7,>=2.5 in /usr/local/lib/python3.6/dist-packages (from requests->smart-open>=1.2.1->gensim) (2.6)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.6/dist-packages (from requests->smart-open>=1.2.1->gensim) (2018.8.24)
Requirement already satisfied: python-dateutil<3.0.0,>=2.1; python_version >= "2.7" in /usr/local/lib/python3.6/dist-packages (from boto3<1.13.0,>=1.12.23->boto3->smart-open>=1.2.1->gensim) (2.5.3)
Requirement already satisfied: docutils>=0.10 in /usr/local/lib/python3.6/dist-packages (from boto3<1.13.0,>=1.12.23->boto3->smart-open>=1.2.1->gensim) (0.14)
```

In [0]:

```
from sklearn.model_selection import train_test_split
X_tra, X_tes, y_train, y_test = train_test_split(cleaned_data['CleanedText'], cleaned_data['Score'].values, test_size=0.3, shuffle=False)
```

In [0]:

```
sent_of_train=[]
for sent in X_tra:
    sent_of_train.append(sent.split())

#sent_of_train

sent_of_test=[]
for sent in X_tes:
    sent_of_test.append(sent.split())
```

In [80]:

```
from gensim.models import Word2Vec
w2v_model=Word2Vec(sent_of_train,min_count=3,size=50, workers=4)# words which occurs 3 times; 50 dimensions

w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 3 times ",len(w2v_words))

from gensim.models import Word2Vec
w2v_model=Word2Vec(sent_of_test,min_count=3,size=50, workers=4)# words which occurs 3 times; 50 dimensions

w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 3 times ",len(w2v_words))
```

```
number of words that occurred minimum 3 times 14278
number of words that occurred minimum 3 times 9641
```

In [81]:

```
# compute average word2vec for each review for X_train .
train_vectors = []
for sent in sent_of_train:
    sent_vec = np.zeros(50)
    cnt_words = 0;
    for word in sent:
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
```

```

        cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    train_vectors.append(sent_vec)

# compute average word2vec for each review for X_test .
test_vectors = []
for sent in sent_of_test:
    sent_vec = np.zeros(50)
    cnt_words = 0;
    for word in sent:
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    test_vectors.append(sent_vec)

X_train = np.asarray(train_vectors)
X_test = np.asarray(test_vectors)

X_train.shape, X_test.shape

```

Out[81]:

```
((70000, 50), (30000, 50))
```

In [82]:

```

# Data-preprocessing: Standardizing the data

sc = StandardScaler(with_mean = False)
X_train1 = sc.fit_transform(X_train)
X_test1 = sc.transform(X_test)

tscv = TimeSeriesSplit(n_splits=4)
for train, cv in tscv.split(X_train1):
    print(X_train1[train].shape, X_train1[cv].shape)

```

```

(14000, 50) (14000, 50)
(28000, 50) (14000, 50)
(42000, 50) (14000, 50)
(56000, 50) (14000, 50)

```

In [107]:

```

#GridSearch with default l2 norm
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import TimeSeriesSplit
from sklearn.metrics import make_scorer
from sklearn.metrics import f1_score

f1 = make_scorer(f1_score, pos_label='positive')
tscv = TimeSeriesSplit(n_splits=3)
param_grid = {'C':[10**-4, 10**-2, 10**0, 10**2, 10**4]}
log = LogisticRegression(class_weight='balanced')
gsv = GridSearchCV(log, param_grid, cv=tscv, scoring=f1)
gsv.fit(X_train1, y_train)

print("Best HyperParameter: ", gsv.best_params_)
print("Best f1: %.2f%%"%(gsv.best_score_*100))

```

```

Best HyperParameter:  {'C': 1}
Best f1: 87.16%

```

In [110]:

```
print(gsv.best_estimator_)
```

```

print("Best HyperParameter: ",gsv.best_params_)
print("Best f1: %.2f%%"%(gsv.best_score_*100))
import matplotlib.pyplot as plt
x=[]
y=[]
for a in gsv.grid_scores_:
    x.append(a[0]['C'])
    y.append(a[1])
plt.xlim(-10,100)
plt.ylim(0.8,1)
plt.xlabel(r"$\ C = 1/Lambda $",fontsize=15)
plt.ylabel("f1")
plt.title(r'F1 v/s $\ C$')
plt.plot(x,y)
plt.show()


```

```

LogisticRegression(C=1, class_weight='balanced', dual=False,
                    fit_intercept=True, intercept_scaling=1, max_iter=100,
                    multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
                    solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
Best HyperParameter: {'C': 1}
Best f1: 87.16%

```

/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_search.py:762: DeprecationWarning: The grid_scores_ attribute was deprecated in version 0.18 in favor of the more elaborate cv_results_ attribute. The grid_scores_ attribute will not be available from 0.20
DeprecationWarning)



In [111]:

```

log = LogisticRegression(class_weight='balanced',C=1000)
log.fit(X_train1,y_train)

pred = (model.predict(X_train1))
print (classification_report(y_train,pred))
result = confusion_matrix(y_train,pred)
print("Confusion Matrix of test set:\n [ [TN  FP]\n [FN  TP] ]\n")
print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')

print("f1 on test set: %0.3f%%"%(f1_score(y_train, pred,pos_label='positive')*100))
print("Non Zero weights:",np.count_nonzero(clf.coef_))


```

	precision	recall	f1-score	support
negative	0.75	0.17	0.28	8172
positive	0.90	0.99	0.94	61828
avg / total	0.88	0.90	0.87	70000

Confusion Matrix of test set:

```
[ [TN  FP]
  [FN  TP] ]
```

```
[[ 1427  6745]
 [  487 61341]]
f1 on test set: 94.433%
Non Zero weights: 50
```



In [112]:

```
clf = LogisticRegression(C=100, penalty='l2')
model=clf.fit(X_train, y_train)
print("Non Zero weights:", np.count_nonzero(clf.coef_))
```

Non Zero weights: 50

In [113]:

```
clf = LogisticRegression(C=100, penalty='l1')
model=clf.fit(X_train, y_train)
print("Non Zero weights:", np.count_nonzero(clf.coef_))
```

Non Zero weights: 50

In [114]:

```
clf = LogisticRegression(C=0.01, penalty='l2')
model=clf.fit(X_train, y_train)
print("Non Zero weights:", np.count_nonzero(clf.coef_))
```

Non Zero weights: 50

In [115]:

```
clf = LogisticRegression(C=0.01, penalty='l1')
model=clf.fit(X_train, y_train)
print("Non Zero weights:", np.count_nonzero(clf.coef_))
```

Non Zero weights: 29

In [0]:

In [116]:

```
from sklearn.metrics import classification_report
```

```

from sklearn.metrics import confusion_matrix
import seaborn as sns

pred = (model.predict(X_test1))
print (classification_report(y_test,pred))
result = confusion_matrix(y_test,pred)
print("Confusion Matrix of test set:\n [ [TN  FP]\n [FN  TP] ]\n")
print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')

print("f1 on test set: %0.3f%%"%(f1_score(y_test, pred,pos_label='positive')*100))

```

	precision	recall	f1-score	support
negative	0.67	0.36	0.47	4098
positive	0.91	0.97	0.94	25902
avg / total	0.87	0.89	0.87	30000

Confusion Matrix of test set:


```

[ [TN  FP]
  [FN  TP] ]

[[ 1478  2620]
 [   712 25190]]

```

f1 on test set: 93.797%



Random search

In [117]:

```

from sklearn.model_selection import RandomizedSearchCV

f1 = make_scorer(f1_score, pos_label='positive')
tscv = TimeSeriesSplit(n_splits=3)
#clf = LogisticRegression(class_weight='balanced',penalty='l2')
#params we need to try on classifier
param_grid = {'C':[10**-4, 10**-2, 10**0, 10**2, 10**4]}

rsv = RandomizedSearchCV(LogisticRegression(class_weight='balanced',penalty='l2'),
param_grid,n_iter=5,cv=tscv,scoring=f1)
rsv.fit(X_train1,y_train)

#savetofile(gsv,"Log Reg/gsv_uni")
print("Best HyperParameter: ",rsv.best_params_)
print("Best F1 score: %0.2f%%"%(rsv.best_score_*100))

```

Best HyperParameter: {'C': 1}

Best F1 score: 87.16%

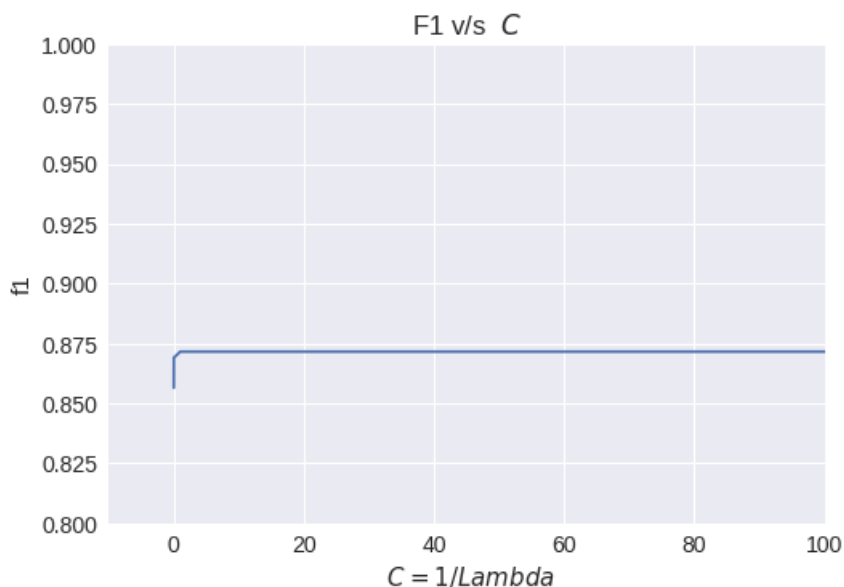
In [118]:

```
In [110]:
```

```
print(rsv.best_estimator_)
print("Best HyperParameter: ",rsv.best_params_)
print("Best f1: %.2f%%"%(rsv.best_score_*100))
import matplotlib.pyplot as plt
x=[]
y=[]
for a in rsv.grid_scores_:
    x.append(a[0]['C'])
    y.append(a[1])
plt.xlim(-10,100)
plt.ylim(0.8,1)
plt.xlabel(r"$\ C = 1/Lambda $",fontsize=15)
plt.ylabel("f1")
plt.title(r'F1 v/s $\ C$')
plt.plot(x,y)
plt.show()
```

```
LogisticRegression(C=1, class_weight='balanced', dual=False,
                    fit_intercept=True, intercept_scaling=1, max_iter=100,
                    multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
                    solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
Best HyperParameter: {'C': 1}
Best f1: 87.16%
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/model_selection/_search.py:762: DeprecationWarning:
The grid_scores_ attribute was deprecated in version 0.18 in favor of the more elaborate
cv_results_ attribute. The grid_scores_ attribute will not be available from 0.20
DeprecationWarning)
```



```
In [123]:
```

```
clf = LogisticRegression(C= 1, penalty= 'l2')
clf.fit(X_train1,y_train)
y_pred = clf.predict(X_train1)

print("Non Zero weights:",np.count_nonzero(clf.coef_))
print("F1-Score on test set: %0.3f%%"%(f1_score(y_train, y_pred, pos_label='positive')*100))
print (classification_report(y_train,y_pred))
result = confusion_matrix(y_train,y_pred)
print("Confusion Matrix of test set:\n [ [TN  FP]\n [FN TP] ]\n")
print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')
```

```
Non Zero weights: 50
F1-Score on test set: 94.659%
      precision    recall  f1-score   support

negative    0.67      0.32      0.43      8172
positive    0.92      0.98      0.95     61828
```

avg / total 0.89 0.90 0.89 70000


Confusion Matrix of test set:

```
[ [TN  FP]
  [FN  TP] ]
```

```
[[ 2592  5580]
 [ 1256 60572]]
```

Out[123]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f88bd747978>



In [125]:

```
pred = (model.predict(X_test1))
print (classification_report(y_test,pred))
result = confusion_matrix(y_test,pred)
print("Confusion Matrix of test set:\n [ [TN  FP]\n [FN  TP] ]\n")
print(result)
sns.set(font_scale=1.4)#for label size
sns.heatmap(result, annot=True,annot_kws={"size": 16}, fmt='g')

print("f1 on test set: %0.3f%%"%(f1_score(y_test, pred,pos_label='positive')*100))
```


	precision	recall	f1-score	support
negative	0.67	0.36	0.47	4098
positive	0.91	0.97	0.94	25902
avg / total	0.87	0.89	0.87	30000

Confusion Matrix of test set:

```
[ [TN  FP]
  [FN  TP] ]
```


```
[[ 1478  2620]
 [  712 25190]]
```

f1 on test set: 93.797%





In [0]:



LR with Grid search CV

- 1. Hyper parameter is 1
- 1. F1 score of Train data = 94.4%
- 1. F1 score of Test data
- 93.7%

LR with Random search CV

- Hyper parameter is 1
- F1 score of Train data = 94.6%
- F1 score of Test data
- 93.7%